

Assignment10_3

November 7, 2021

```
[10]: import tensorflow.compat.v1 as tf
tf.disable_v2_behavior()
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
import numpy as np
import matplotlib.pyplot as plt
from pathlib import Path
from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense
import os
from contextlib import redirect_stdout
import time
start_time = time.time()
from keras.layers import LSTM
from keras.datasets import imdb
from keras.preprocessing import sequence
```

```
[2]: imdb_dir = Path('/home/jovyan/dsc650/data/external/imdb/aclImdb/')
test_dir = os.path.join(imdb_dir, 'test')
train_dir = os.path.join(imdb_dir, 'train')
results_dir = Path('results').joinpath('model_1')
results_dir.mkdir(parents=True, exist_ok=True)
```

```
[3]: max_features = 10000
maxlen = 500
batch_size = 32
max_words = 1000
training_samples = 200
validation_samples = 10000
```

```
[4]: labels = []
texts = []
for label_type in ['neg', 'pos']:
    dir_name = os.path.join(test_dir, label_type)
    for fname in sorted(os.listdir(dir_name)):
        if fname[-4:] == '.txt':
            f = open(os.path.join(dir_name, fname), encoding="utf8")
```

```

        texts.append(f.read())
    f.close()
    if label_type == 'neg':
        labels.append(0)
    else:
        labels.append(1)

```

```

[5]: tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)
print('Loading data... ')
word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))
data = pad_sequences(sequences, maxlen=maxlen)
labels = np.asarray(labels)
print('Shape of data tensor:', data.shape)
print('Shape of label tensor:', labels.shape)
indices = np.arange(data.shape[0])
np.random.shuffle(indices)
data = data[indices]
labels = labels[indices]

```

Loading data...
 Found 87393 unique tokens.
 Shape of data tensor: (25000, 500)
 Shape of label tensor: (25000,)

```

[6]: #x_train
input_train = data[:training_samples]
#y_train
y_train = labels[:training_samples]
#x_val
input_test = data[training_samples: training_samples + validation_samples]
#y_val
y_test = labels[training_samples: training_samples + validation_samples]
print('input_train shape:', input_train.shape)
print('input_test shape:', input_test.shape)

```

input_train shape: (200, 500)
 input_test shape: (10000, 500)

```

[7]: # from page 205 - Listing 6.27
model = Sequential()
model.add(Embedding(max_features, 32))
model.add(LSTM(32))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])

```

```

history=model.fit(input_train, y_train, epochs=10,
    ↪batch_size=32,validation_data=(input_test, y_test))
result_model_file = results_dir.joinpath('pre_trained_glove_model_LSTM.h5')
model.save_weights(result_model_file)

```

```

WARNING:tensorflow:From /opt/conda/lib/python3.8/site-
packages/tensorflow/python/keras/initializers/initializers_v1.py:58: calling
RandomUniform.__init__ (from tensorflow.python.ops.init_ops) with dtype is
deprecated and will be removed in a future version.
Instructions for updating:
Call initializer instance with the dtype argument instead of passing it to the
constructor
Train on 200 samples, validate on 10000 samples
Epoch 1/10
200/200 [=====] - ETA: 0s - loss: 0.6941 - acc: 0.5000

/opt/conda/lib/python3.8/site-
packages/tensorflow/python/keras/engine/training.py:2325: UserWarning:
`Model.state_updates` will be removed in a future version. This property should
not be used in TensorFlow 2.0, as `updates` are applied automatically.
  warnings.warn("`Model.state_updates` will be removed in a future version. '



200/200 [=====] - 23s 114ms/sample - loss: 0.6941 -
acc: 0.5000 - val_loss: 0.6921 - val_acc: 0.5051
Epoch 2/10
200/200 [=====] - 21s 103ms/sample - loss: 0.6874 -
acc: 0.5450 - val_loss: 0.6907 - val_acc: 0.5545
Epoch 3/10
200/200 [=====] - 21s 106ms/sample - loss: 0.6782 -
acc: 0.7900 - val_loss: 0.6876 - val_acc: 0.6027
Epoch 4/10
200/200 [=====] - 21s 105ms/sample - loss: 0.6579 -
acc: 0.8400 - val_loss: 0.6793 - val_acc: 0.6118
Epoch 5/10
200/200 [=====] - 21s 105ms/sample - loss: 0.6568 -
acc: 0.8250 - val_loss: 0.6476 - val_acc: 0.6766
Epoch 6/10
200/200 [=====] - 21s 105ms/sample - loss: 0.5341 -
acc: 0.8700 - val_loss: 0.6322 - val_acc: 0.6365
Epoch 7/10
200/200 [=====] - 21s 106ms/sample - loss: 0.4548 -
acc: 0.8700 - val_loss: 0.6299 - val_acc: 0.6561
Epoch 8/10
200/200 [=====] - 21s 106ms/sample - loss: 0.3917 -
acc: 0.9400 - val_loss: 1.0575 - val_acc: 0.5094
Epoch 9/10
200/200 [=====] - 21s 106ms/sample - loss: 0.3532 -
acc: 0.9200 - val_loss: 0.5785 - val_acc: 0.6949

```

```
Epoch 10/10
200/200 [=====] - 21s 104ms/sample - loss: 0.2930 -
acc: 0.9050 - val_loss: 0.5757 - val_acc: 0.6961
```

```
[8]: # Save the summary to file
summary_file = results_dir.joinpath('Assignment_10.3_ModelSummary.txt')
with open(summary_file, 'w') as f:
    with redirect_stdout(f):
        model.summary()

[9]: # Plots
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
img_file = results_dir.joinpath('Assignment_10.3_Model Accuracy Validation.png')
plt.savefig(img_file)
plt.show()
```



[]: